

# HW3 Report

111065541 張騰午

## Implement

### Hw3-1

做 Floyd-Warshall，時間複雜度  $O(n^3)$ ，不過在  $i, j$  兩個迴圈增加 `omp` 指令來平行化。

```
for (int k = 0; k < vertex; ++k){
    #pragma omp parallel for schedule(guided, 1) collapse(2)
    for (int i = 0; i < vertex; ++i){
        for (int j = 0; j < vertex; ++j){
            if (Dist[i][j] > Dist[i][k] + Dist[k][j] && Dist[i][k] != INF)
                Dist[i][j] = Dist[i][k] + Dist[k][j];
        }
    }
}
```

另外在 initial 部分使用 `fill` 一次性填入 `INF` 並用 `omp` 指令來增加速度


```
#pragma omp parallel for schedule(guided, 1)
for (int i = 0; i < vertex; ++i) {
    fill(Dist[i].begin(), Dist[i].end(), INF);
    Dist[i][i] = 0;
}
```

### Hw3-2

使用了助教提供範例 blocked-Floyd-Warshall 演算法，input 時修改使用 `int *fpt = (int*)mmap(NULL, 2*sizeof(int), PROT_READ, MAP_PRIVATE, file, 0);` 如此直接將所有資料一次性讀取，之後再寫入  $V \times V$  的 `Dist` 中，如此可以減少 I/O 時間。

而演算法計算方面思路如下：

blocking factor 64，而一個 GPU 的 block 為  $32 \text{ threads} \times 32 \text{ threads}$ ，因此整個演算法把矩陣分成數個  $64 \times 64$  block，每個 block 使用一個 GPU  $32 \times 32$  的 block 下去算要計算 4 次，一個 thread 在一個  $B \times B$  block 要算 4 個點，如下圖紅色標示處，4 個點距離左右、上下皆為 32，



三個 Phase 分別用一個  $32 \times 32$  block\_dim threads、兩行  $\text{ceil}(V/B)$  個  $32 \times 32$  block\_dim threads 以及  $\text{ceil}(V/B) \times \text{ceil}(V/B)$  個  $32 \times 32$  block\_dim threads 平面進行計算，宣告方式如下：

```
int blocks = ceil(V, B);
dim3 block_dim(BLOCK_SIZE, BLOCK_SIZE); // BLOCK_SIZE=32
dim3 grid_dim(blocks, blocks);
dim3 grid_dim2(blocks, 2);

for (int r = 0; r < round; ++r) {
    // phase 1
    Phase1<<<1, block_dim>>>(dst, r, V);
    // phase 2
    Phase2<<<grid_dim2, block_dim>>>(dst, r, V);
    // phase 3
    Phase3<<<grid_dim, block_dim>>>(dst, r, V);
    //Phase3<<<grid_dim, block_dim>>>(dst, r, V);
}
```

Phase 中需要計算的點座標如下(用一維方式表示)，

Phase 1(僅 pivot block)

```
__global__ void Phase1(int *dst, int Round, int V){
    int i = threadIdx.y;
    int j = threadIdx.x;
    int i_B = i + BLOCK_SIZE;
    int j_B = j + BLOCK_SIZE;
    //B為64，一次處理一個64*64，但block dim僅為32*32，因此一個大block分成四個小block，一個thread要計算四個小block四個點
    // 1 2
    // 3 4
    int offset = Round * B * (V+1);
    int blk_pt1 = offset + i * V + j;
    int blk_pt2 = offset + i * V + j_B;
    int blk_pt3 = offset + i_B * V + j;
    int blk_pt4 = offset + i_B * V + j_B;
```

Phase 2(下左，依據 blockIdx.y 的值判斷計算 pivot row 的 block 或是 pivot column 的 block)

Phase 3(下右，使用同 row 還有同 column 的 pivot row/column block 來計算值)

```

int blk_pt1 = offset + i * V + j;
int blk_pt2 = offset + i * V + j_B;
int blk_pt3 = offset + i_B * V + j;
int blk_pt4 = offset + i_B * V + j_B;
//for same col block
if(blockIdx.y == 0) {
    offset_rc = blockIdx.x * B * V + Round * B;
    pivot_rc_blk1 = offset_rc + i * V + j;
    pivot_rc_blk2 = offset_rc + i * V + j_B;
    pivot_rc_blk3 = offset_rc + i_B * V + j;
    pivot_rc_blk4 = offset_rc + i_B * V + j_B;
} else { //for same row block
    offset_rc = Round * B * V + blockIdx.x * B;
    pivot_rc_blk1 = offset_rc + i * V + j;
    pivot_rc_blk2 = offset_rc + i * V + j_B;
    pivot_rc_blk3 = offset_rc + i_B * V + j;
    pivot_rc_blk4 = offset_rc + i_B * V + j_B;
}

int offset_ = blockIdx.y * B * V + blockIdx.x * B;
int blk_pt1 = offset_ + i * V + j;
int blk_pt2 = offset_ + i * V + j_B;
int blk_pt3 = offset_ + i_B * V + j;
int blk_pt4 = offset_ + i_B * V + j_B;
//same row
int offset_r = blockIdx.y * B * V + Round * B;
int row_blk1 = offset_r + i * V + j;
int row_blk2 = offset_r + i * V + j_B;
int row_blk3 = offset_r + i_B * V + j;
int row_blk4 = offset_r + i_B * V + j_B;
//same col
int offset_c = Round * B * V + blockIdx.x * B;
int col_blk1 = offset_c + i * V + j;
int col_blk2 = offset_c + i * V + j_B;
int col_blk3 = offset_c + i_B * V + j;
int col_blk4 = offset_c + i_B * V + j_B;

```

三個 phase 結束後先在 gpu 上用 devicetodevice 的 CudaMemcpy 把本來  $V \times V$  存在 dst 的資料複製到另一個  $\text{vertex} \times \text{vertex}$  大小的 dst\_，

```

cudaMalloc(&dst_, vertex*vertex*sizeof(int));
for(int i = 0; i < vertex; ++i) {
    cudaMemcpy(dst_ + i*vertex, dst + i*V, sizeof(int)*vertex, cudaMemcpyDeviceToDevice);
}
cudaMemcpy(Dist, dst_, size_vertex, cudaMemcpyDeviceToHost);

```

之後 output 時就可以直接 `fwrite(Dist, sizeof(int), vertex*vertex, outfile);` 一次性把  $\text{vertex} \times \text{vertex}$  量的資料全部寫入 output file 中減少 I/O 時間。


用 `__shared__ int share[B][B];` 創建 share memory，來讀取後續比較還有計算需要的值，每個 thread 將自己負責的值存入 share memory，並進行後續更新值的計算，存取修改 share memory 後要 `__syncthreads();` 進行同步，最後把 share memory 中的值存回 dst 中。

另外在進行比較計算時，發現使用 `__device__` 函數再調用 cuda 內建函數 min，讓計算在 GPU 上進行會比直接用 if 判斷式更快。

```
__device__ int min_(int a, int b) {return min(a, b);}
```

### Hw3-3

3-2 的程式大致不變，查詢相關資料並和同學討論後只在 phase 3 使用兩 GPU 平分執行，因為有了 start\_point，phase 3 計算座標時要多加入 y 軸方面 offset。平分方式如下圖，整個矩陣本來能分成  $\text{round} \times \text{round}$  個  $B \times B$  block，在 y 軸上分成上下兩部份給兩個 GPU 做。



程式在執行三個 phase 時用 `#pragma omp parallel num_threads(2)` 來執行兩 gpu，接著取得 `thread_id` 以及 `neighbor thread id`，並計算雙方需要做的數量還有起始點，若是 `round` 無法被整除，則將多的一個 `round` 加在 `thread_id=0` 的那邊

```
unsigned int thread_id = omp_get_thread_num();
cudaSetDevice(thread_id);
// neighbor
unsigned int neighbor_thread_id;
if(thread_id){
    neighbor_thread_id = 0;
}else{
    neighbor_thread_id = 1;
}

unsigned int rounds_to_do = round / 2;
unsigned int rounds_to_do_neighbor = round / 2;
if (thread_id == 0) {
    rounds_to_do += round % 2;
}else{rounds_to_do_neighbor += round % 2;}

unsigned int start_point;
if (thread_id) {
    start_point = rounds_to_do_neighbor;
} else {
    start_point = 0;
}
```

進入迴圈後每次都得進行迴圈 `r` 對應的那一行 `block` 數據交換(pivot row)，通過 `#pragma omp barrier` 來確保兩邊都已經完成數據操作。

```
for (int r = 0; r < round; ++r) {
    if (r >= start_point && r < (start_point + rounds_to_do)) {
        cudaMemcpy(dst[neighbor_thread_id] + r * B * V, dst[thread_id] + r * B * V, B * V * sizeof(int), cudaMemcpyDefault)
    }
    #pragma omp barrier
}
```

## Experiment & Analysis

### i、 Methodology

#### a. System Spec

在課程提供的hades上進行測試

NVIDIA-SMI 384.81				Driver Version: 384.81			
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M.
0	GeForce GTX 1080	On	00000000:4B:00:0	Off			N/A
32%	47C	P8	14W / 216W	2MiB / 8114MiB	0%	Default	
1	GeForce GTX 1080	On	00000000:4D:00:0	Off			N/A
10%	47C	P8	18W / 216W	11MiB / 8113MiB	0%	Default	

## b. Blocking Factor (hw3-2)

### GOPS

c21.1 進行 --metrics inst\_integer 計算再除以執行時間

B=64

```

[pp23s7]@hades01 hw3-21$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics inst_integer ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222856== NVPROF is profiling process 222856, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222856== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222856== Profiling result:
==222856== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
79              inst_integer      Integer Instructions     98244608  98244608  98244608
Kernel: Phase1(int*, int, int)
79              inst_integer      Integer Instructions     610304    610304    610304
Kernel: Phase3(int*, int, int)
79              inst_integer      Integer Instructions     2859898880 2859898880 2859898880

```

```

[pp23s7]@hades01 hw3-21$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222921== NVPROF is profiling process 222921, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222921== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64.out
==222921== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 73.15%      128.03ms      79      1.6207ms      1.5063ms      1.6559ms      Phase3(int*, int, int)
9.10%      15.930ms      1      15.930ms      15.930ms      15.930ms      [CUDA memcpy HtoD]
8.94%      15.641ms      1      15.641ms      15.641ms      15.641ms      [CUDA memcpy DtoH]
7.62%      13.334ms      79      168.78us      152.77us      178.27us      Phase2(int*, int, int)
1.19%      2.0850ms      79      26.392us      24.128us      27.169us      Phase1(int*, int, int)
API calls: 62.23%      174.58ms      2      87.288ms      15.953ms      158.62ms      cudaMemcpy
79.41%      104.94ms      1      104.94ms      104.94ms      104.94ms      cudaHostRegister
0.22%      624.42us      237      2.6340us      2.0330us      84.972us      cudaLaunchKernel
0.05%      147.82us      1      147.82us      147.82us      147.82us      cudaFree
0.05%      143.90us      101      1.4240us      103ns      60.857us      cuDeviceGetAttribute
0.03%      93.803us      1      93.803us      93.803us      93.803us      cudaMalloc
0.00%      10.844us      1      10.844us      10.844us      10.844us      cuDeviceGetName
0.00%      7.8010us      1      7.8010us      7.8010us      7.8010us      cuDeviceGetPCIBusId
0.00%      1.1290us      3      376ns      167ns      738ns      cuDeviceGetCount
0.00%      600ns      1      600ns      600ns      600ns      cuModuleGetLoadingMode
0.00%      493ns      2      246ns      129ns      364ns      cuDeviceGet
0.00%      267ns      1      267ns      267ns      267ns      cuDeviceTotalMem
0.00%      190ns      1      190ns      190ns      190ns      cuDeviceGetUuid

```

B=32

```

[pp23s7]@hades01 hw3-21$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics inst_integer ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==123134== NVPROF is profiling process 123134, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==123134== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==123134== Profiling result:
==123134== Metric result:
Invocations      Metric Name      Metric Description      Min      Max      Avg
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
157          inst_integer      Integer Instructions     26118656  26118656  26118656
Kernel: Phase1(int*, int, int)
157          inst_integer      Integer Instructions     78848     78848     78848
Kernel: Phase3(int*, int, int)
157          inst_integer      Integer Instructions     1563894272 1563894272 1563894272

```

```

[pp23s7]@hades01 hw3-21$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==222785== NVPROF is profiling process 222785, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==222785== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32.out
==222785== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 78.39%      158.09ms      157      1.0070ms      984.43us      1.0314ms      Phase3(int*, int, int)
7.79%      15.701ms      1      15.701ms      15.701ms      15.701ms      [CUDA memcpy HtoD]
7.66%      15.452ms      1      15.452ms      15.452ms      15.452ms      [CUDA memcpy DtoH]
5.48%      11.043ms      157      70.334us      65.345us      73.703us      Phase2(int*, int, int)
0.68%      1.3765ms      157      8.7670us      8.0960us      9.4080us      Phase1(int*, int, int)
API calls: 65.80%      200.66ms      2      100.33ms      15.724ms      184.94ms      cudaMemcpy
0.41%      1.2393ms      471      2.6310us      2.0380us      82.213us      cudaHostRegister
33.67%      102.69ms      1      102.69ms      102.69ms      102.69ms      cudaLaunchKernel
0.05%      140.89us      101      1.3940us      102ns      61.210us      cuDeviceGetAttribute
0.04%      128.59us      1      128.59us      128.59us      128.59us      cudaFree
0.03%      86.052us      1      86.052us      86.052us      86.052us      cudaMalloc
0.00%      11.598us      1      11.598us      11.598us      11.598us      cuDeviceGetName
0.00%      8.4310us      1      8.4310us      8.4310us      8.4310us      cuDeviceGetPCIBusId
0.00%      1.1830us      3      394ns      180ns      817ns      cuDeviceGetCount
0.00%      704ns      2      352ns      124ns      580ns      cuDeviceGet
0.00%      312ns      1      312ns      312ns      312ns      cuModuleGetLoadingMode
0.00%      311ns      1      311ns      311ns      311ns      cuDeviceTotalMem
0.00%      283ns      1      283ns      283ns      283ns      cuDeviceGetUuid

```



B=16

```

[pp23s71@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics inst_integer ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==158005== NVPROF is profiling process 158005, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==158005== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==158005== Profiling result:
==158005== Metric result:
Invocations                                     Metric Name                                     Metric Description                                     Min                                     Max                                     Avg
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
313                                             inst_integer                                     Integer Instructions                                     7308416                                7308416                                7308416
Kernel: Phase1(int*, int, int)
313                                             inst_integer                                     Integer Instructions                                     10496                                 10496                                 10496
Kernel: Phase3(int*, int, int)
313                                             inst_integer                                     Integer Instructions                                     866052224                             866052224                             866052224

```

```

[pp23s71@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==223809== NVPROF is profiling process 223809, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==223809== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16.out
==223809== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:
91.29%    422.51ms    313      1.3499ms  1.2889ms  1.4326ms  Phase3(int*, int, int)
3.41%     15.770ms    1        15.770ms  15.770ms  15.770ms  [CUDA memcpy HtoD]
3.31%     15.309ms    1        15.309ms  15.309ms  15.309ms  [CUDA memcpy DtoH]
1.67%     7.7137ms    313      24.644us  22.528us  27.137us  Phase2(int*, int, int)
0.33%     1.5164ms    313      4.8440us  4.4160us  5.3760us  Phase1(int*, int, int)
API calls:
80.50%    460.93ms    2        230.47ms  15.854ms  445.08ms  cudaMemcpy
19.03%    108.99ms    1        108.99ms  108.99ms  108.99ms  cudaHostRegister
0.40%     2.2719ms    939      2.4190us  2.0280us  84.584us  cudaLaunchKernel
0.03%     158.05us    1        158.05us  158.05us  158.05us  cudaFree
0.02%     138.60us    101      1.3720us  105ns     60.827us  cuDeviceGetAttribute
0.02%     91.753us    1        91.753us  91.753us  91.753us  cudaMalloc
0.00%     9.2720us    1        9.2720us  9.2720us  9.2720us  cuDeviceGetName
0.00%     8.3130us    1        8.3130us  8.3130us  8.3130us  cuDeviceGetPCIBusId
0.00%     1.1020us    3        367ns     149ns     759ns     cuDeviceGetCount
0.00%     567ns       2        283ns     125ns     442ns     cuDeviceGet
0.00%     547ns       1        547ns     547ns     547ns     cuModuleGetLoadingMode
0.00%     338ns       1        338ns     338ns     338ns     cuDeviceTotalMem
0.00%     196ns       1        196ns     196ns     196ns     cuDeviceGetUuid

```

B=8

```


[pp23s71@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics inst_integer ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==223926== NVPROF is profiling process 223926, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==223926== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==223926== Profiling result:
==223926== Metric result:
Invocations                                     Metric Name                                     Metric Description                                     Min                                     Max                                     Avg
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
625                                             inst_integer                                     Integer Instructions                                     2076704                                2076704                                2076704
Kernel: Phase1(int*, int, int)
625                                             inst_integer                                     Integer Instructions                                     1344                                 1344                                 1344
Kernel: Phase3(int*, int, int)
625                                             inst_integer                                     Integer Instructions                                     566971424                             566971424                             566971424

```

```

[pp23s71@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==158503== NVPROF is profiling process 158503, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==158503== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8.out
==158503== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:
97.76%    1.75083s    625      2.8013ms  2.5571ms  2.9828ms  Phase3(int*, int, int)
0.87%     15.496ms    1        15.496ms  15.496ms  15.496ms  [CUDA memcpy HtoD]
0.85%     15.214ms    1        15.214ms  15.214ms  15.214ms  [CUDA memcpy DtoH]
0.41%     7.2831ms    625      11.652us  10.880us  12.928us  Phase2(int*, int, int)
0.11%     2.0435ms    625      3.2690us  3.0720us  3.9360us  Phase1(int*, int, int)
API calls:
52.10%    985.55ms    2        492.77ms  15.534ms  970.02ms  cudaMemcpy
42.60%    805.82ms    1875     429.77us  1.6900us  2.9847ms  cudaLaunchKernel
5.28%     99.946ms    1        99.946ms  99.946ms  99.946ms  cudaHostRegister
0.01%     161.96us    1        161.96us  161.96us  161.96us  cudaFree
0.01%     120.33us    101      1.1910us  91ns      52.429us  cuDeviceGetAttribute
0.00%     71.254us    1        71.254us  71.254us  71.254us  cudaMalloc
0.00%     8.1560us    1        8.1560us  8.1560us  8.1560us  cuDeviceGetName
0.00%     6.3610us    1        6.3610us  6.3610us  6.3610us  cuDeviceGetPCIBusId
0.00%     916ns       3        305ns     150ns     598ns     cuDeviceGetCount
0.00%     648ns       2        324ns     115ns     533ns     cuDeviceGet
0.00%     454ns       1        454ns     454ns     454ns     cuModuleGetLoadingMode
0.00%     314ns       1        314ns     314ns     314ns     cuDeviceTotalMem
0.00%     174ns       1        174ns     174ns     174ns     cuDeviceGetUuid

```



從圖中可發現隨著 blocking factor 增加，每秒的 operation 也在增加，推測這是因為並行化程度增加，同時進行的計算操作變多的緣故。

## global/shared memory bandwidth

### B64\*64

```
[pp23s1@shades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics shared_store_throughput,shared_load_throughput,gld_throughput,gst_throughput ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64_64.out
srun: job 803094 queued and waiting for resources
srun: job 803094 has been allocated resources
==226715== NVPROF is profiling process 226715, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64_64.out
==226715== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==226715== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B64_64.out
==226715== Profiling result:
==226715== Metric result:
Invocations
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
79 shared_store_throughput Shared Memory Store Throughput 913.39GB/s 1051.4GB/s 1032.6GB/s
79 shared_load_throughput Shared Memory Load Throughput 2228.1GB/s 2564.8GB/s 2519.0GB/s
79 gld_throughput Global Load Throughput 27.67GB/s 31.88GB/s 31.29GB/s
79 gst_throughput Global Store Throughput 13.83GB/s 15.93GB/s 15.64GB/s
Kernel: Phase1(int*, int, int)
79 shared_store_throughput Shared Memory Store Throughput 37.46GB/s 42.86GB/s 42.30GB/s
79 shared_load_throughput Shared Memory Load Throughput 111.24GB/s 127.29GB/s 125.63GB/s
79 gld_throughput Global Load Throughput 590.20MB/s 675.35MB/s 666.53MB/s
79 gst_throughput Global Store Throughput 590.20MB/s 675.35MB/s 666.53MB/s
Kernel: Phase3(int*, int, int)
79 shared_store_throughput Shared Memory Store Throughput 249.06GB/s 260.08GB/s 258.23GB/s
79 shared_load_throughput Shared Memory Load Throughput 3051.0GB/s 3186.0GB/s 3163.3GB/s
79 gld_throughput Global Load Throughput 186.80GB/s 195.06GB/s 193.67GB/s
79 gst_throughput Global Store Throughput 62.26GB/s 65.02GB/s 64.55GB/s
```

### B32\*32

```

[pp23s1@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics shared_store_throughput,shared_load_throughput,gld_throughput,gst_throughput ./w3-2 /home/pp23/share/hw3-2/cases/c21.1 B32_32.out
==227347== NVPROF is profiling process 227347, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32_32.out
==227347== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==227347== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B32_32.out
==227347== Profiling result:
==227347== Metric result:
Invocations
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
157 shared_store_throughput Shared Memory Store Throughput 1189.8GB/s 1266.2GB/s 1227.2GB/s
157 shared_load_throughput Shared Memory Load Throughput 2414.6GB/s 2569.5GB/s 2490.6GB/s
157 gld_throughput Global Load Throughput 34.995GB/s 37.240GB/s 36.095GB/s
157 gst_throughput Global Store Throughput 17.497GB/s 18.620GB/s 18.048GB/s
Kernel: Phase1(int*, int, int)
157 shared_store_throughput Shared Memory Store Throughput 30.378GB/s 33.480GB/s 31.874GB/s
157 shared_load_throughput Shared Memory Load Throughput 74.563GB/s 82.178GB/s 78.236GB/s
157 gld_throughput Global Load Throughput 471.31MB/s 519.45MB/s 494.53MB/s
157 gst_throughput Global Store Throughput 471.31MB/s 519.45MB/s 494.53MB/s
Kernel: Phase3(int*, int, int)
157 shared_store_throughput Shared Memory Store Throughput 707.15GB/s 725.08GB/s 717.46GB/s
157 shared_load_throughput Shared Memory Load Throughput 2290.2GB/s 2356.9GB/s 2331.6GB/s
157 gld_throughput Global Load Throughput 265.18GB/s 271.91GB/s 269.05GB/s
157 gst_throughput Global Store Throughput 88.394GB/s 90.635GB/s 89.683GB/s

```

## B16\*16

```

[pp23s1@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics shared_store_throughput,shared_load_throughput,gld_throughput,gst_throughput ./w3-2 /home/pp23/share/hw3-2/cases/c21.1 B16_16.out
==160848== NVPROF is profiling process 160848, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16_16.out
==160848== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==160848== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B16_16.out
==160848== Profiling result:
==160848== Metric result:
Invocations
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
313 shared_store_throughput Shared Memory Store Throughput 782.56GB/s 904.70GB/s 867.49GB/s
313 shared_load_throughput Shared Memory Load Throughput 1521.6GB/s 1759.1GB/s 1686.8GB/s
313 gld_throughput Global Load Throughput 43.476GB/s 50.261GB/s 48.194GB/s
313 gst_throughput Global Store Throughput 43.476GB/s 50.261GB/s 48.194GB/s
Kernel: Phase1(int*, int, int)
313 shared_store_throughput Shared Memory Store Throughput 6.6663GB/s 8.1881GB/s 7.6793GB/s
313 shared_load_throughput Shared Memory Load Throughput 16.078GB/s 19.748GB/s 18.521GB/s
313 gld_throughput Global Load Throughput 200.77MB/s 246.61MB/s 231.28MB/s
313 gst_throughput Global Store Throughput 401.55MB/s 493.21MB/s 462.57MB/s
Kernel: Phase3(int*, int, int)
313 shared_store_throughput Shared Memory Store Throughput 563.57GB/s 603.74GB/s 580.49GB/s
313 shared_load_throughput Shared Memory Load Throughput 986.25GB/s 1056.5GB/s 1015.9GB/s
313 gld_throughput Global Load Throughput 211.34GB/s 226.40GB/s 217.68GB/s
313 gst_throughput Global Store Throughput 140.89GB/s 150.94GB/s 145.12GB/s

```



## B8\*8

```

[pp23s1@hades01 hw3-2]$ srun -p prof -N1 -n1 --gres=gpu:1 nvprof --metrics shared_store_throughput,shared_load_throughput,gld_throughput,gst_throughput ./w3-2 /home/pp23/share/hw3-2/cases/c21.1 B8_8.out
==161430== NVPROF is profiling process 161430, command: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8_8.out
==161430== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==161430== Profiling application: ./hw3-2 /home/pp23/share/hw3-2/cases/c21.1 B8_8.out
==161430== Profiling result:
==161430== Metric result:
Invocations
Device "NVIDIA GeForce GTX 1080 (0)"
Kernel: Phase2(int*, int, int)
625 shared_store_throughput Shared Memory Store Throughput 464.92GB/s 524.96GB/s 495.99GB/s
625 shared_load_throughput Shared Memory Load Throughput 976.32GB/s 1102.4GB/s 1041.6GB/s
625 gld_throughput Global Load Throughput 92.983GB/s 104.99GB/s 99.198GB/s
625 gst_throughput Global Store Throughput 46.492GB/s 52.496GB/s 49.599GB/s
Kernel: Phase1(int*, int, int)
625 shared_store_throughput Shared Memory Store Throughput 1.2895GB/s 1.5778GB/s 1.5216GB/s
625 shared_load_throughput Shared Memory Load Throughput 3.5820GB/s 4.3827GB/s 4.2267GB/s
625 gld_throughput Global Load Throughput 146.72MB/s 179.52MB/s 173.12MB/s
625 gst_throughput Global Store Throughput 146.72MB/s 179.52MB/s 173.12MB/s
Kernel: Phase3(int*, int, int)
625 shared_store_throughput Shared Memory Store Throughput 244.58GB/s 258.12GB/s 250.31GB/s
625 shared_load_throughput Shared Memory Load Throughput 428.02GB/s 451.72GB/s 438.04GB/s
625 gld_throughput Global Load Throughput 183.44GB/s 193.59GB/s 187.73GB/s
625 gst_throughput Global Store Throughput 61.146GB/s 64.531GB/s 62.577GB/s

```



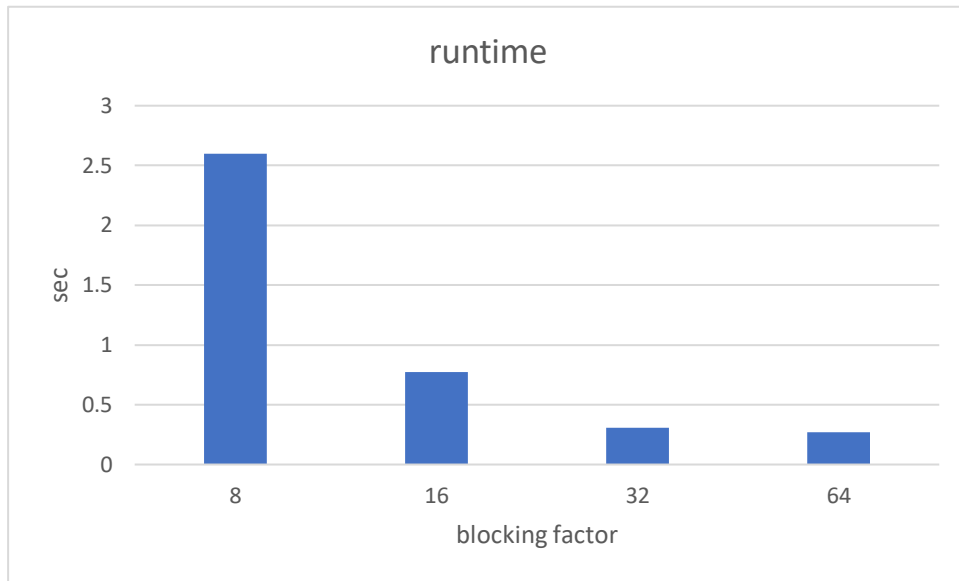


由於程式中對於 share memory 的存取計算，Share memory store 隨著 blocking factor 增加有非常明顯的差異，幅度上千，但 share memory 也有數量限制，因此不會一直上升。global 部份則沒有那麼劇烈的變化。

### c. Optimization (hw3-2)

#### 1. Large blocking factor


c21.1 進行測試四種 blocking factor 下的執行時間



較大的 blocking factor 的確可以明顯降低運行時間

## 2. share memory

使用 share memory 可以達到很好的優化效果，對 c21.1 比較使用 share memory 還有 no share memory，兩者差距是巨量的。



## 3. 比較計算優化

以 p16k1 V=16000 測資，在 block-floyd-warshall 計算時比較使用 \_\_device\_\_ 函數


```
__device__ int min_(int a, int b) {return min(a, b);}
```

、在 GPU 上執行再調用 cuda 內建函數 min，以及直接使用 if

還有在 #define 定義一個 macro：

```
#define min__(a, b) ((a) < (b) ? (a) : (b))
```

三種方式進行比較：



測量方式是使用 `cudaEvent` 計算每部分累計時間。可以清楚看到使用 `__device__` 與 `cuda` 內建 `min` 函數在 GPU 上執行是最快的，而以 `#define` 寫一個比較大小的 `macro` 還有 `if` 則慢很多，差異接近兩倍。

#### 4. IO 優化

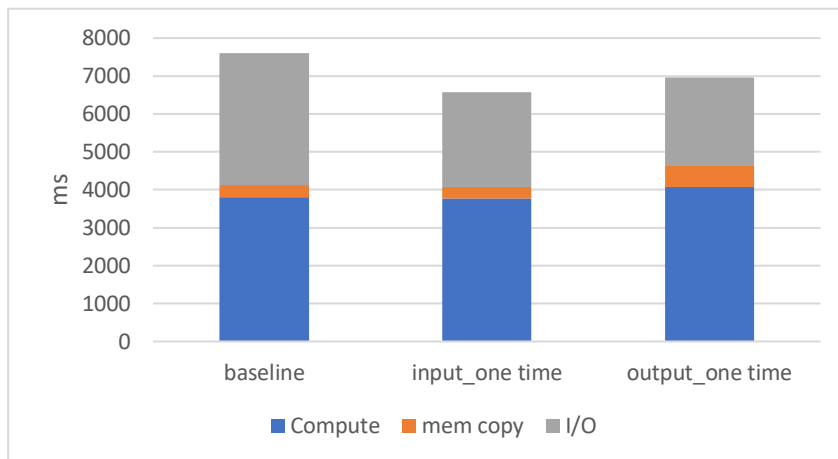
另外進行了 IO 相關優化實驗，同樣使用 p16k1 測資，測量方式是使用 `cudaEvent` 計算每部分累計時間。

Baseline 是使用原來助教範例程式中的 `input` 與 `output` 程式

`Input_one time` 是使用

`int *fpt = (int*)mmap(NULL, 2*sizeof(int), PROT_READ, MAP_PRIVATE, file, 0);` 直接將所有資料一次性讀取，之後再寫入 `V*V` 的 `Dist` 中，

而 `output_one time` 則是在先三個 `phase` 結束後先在 `gpu` 上用 `devicetodevice` 的 `CudaMemcpy` 把本來 `V*V` 存在 `dst` 的資料複製到另一個 `vertex*vertex` 大小的 `dst_`，之後 `output` 時就可以直接 `fwrite(Dist, sizeof(int), vertex*vertex, outfile);` 一次性把 `vertex*vertex` 量的資料全部寫入 `output file` 中減少 I/O 時間。




結果可看出 IO 時間都有降低，尤其 input\_one time 十分顯著，output\_one time 雖然也講了不少 IO 時間，但因為要做更多的 CudaMemcpy 所以 memory copy 時間也有增長。

#### d. Weak scalability (hw3-3)


Weak scalability 是在每個計算單元計算量相同時比較不同數量計算單元能否在差不多時間內完成任務，在此使用數量為 1 與 2 的 GPU 進行測試，直接修改程式內 V 的值測試計算時間。測量方式是使用 cudaEvent 計算每部分累計時間。

兩個 GPU 計算量應為一個的兩倍( $V \times V$  兩倍)，故 V 應為 $\sqrt{2}$ 倍，這裡單 GPU 使用 16000 與其 $\sqrt{2}$ 倍也就是雙 GPU V 為約 22628 比較，單位為 ms。

可看出雙 GPU 仍高出單 GPU 一些時間，執行時間約為單 GPU 1.2~1.3 倍，可發現有較明顯增長的是 I/O 以及 memory copy，I/O 方面已經有使用一次性 mmap 毒入來節省時間，但後續將內容寫入 Dist 應仍會造成時間差異，此外雙 GPU 運行時每個迴圈要進行兩個 GPU 間 CudaMemcpy，這也造成 memory copy 時間明顯增長，而 compute 時間因為有兩個 GPU 分工是增長較不明顯的。



#### e. Time Distribution (hw3-2)



測試六個V值不同的testcase，測量方式是使用cudaEvent計算每部分累計時間。可看出Memory copy時間是佔比最少的，其上漲幅度也很低，主因是memory copy時間主要在執行cudaMemcpy時，此函數主要是在copy  $V \times V$  的資料量，因此只和input問題的節點數量有關。


接著是I/O時間，I/O時間，同樣來自於input問題大小，由於input, output兩個函數要不斷進行fread,fwrite，其成長幅度高於memory copy。

最後是compute time，compute time不只和input節點數量有關，也和input距離矩陣內容有關，其內容會影響計算量，其時間上升幅度也是三者中最大的。

## f. Others


1. 在 large blocking factor 部分可看出 large blocking factor 的確有益於執行時間，這裡進一步觀察 speedup 程度





其實我本來以為 **blocking factor** 的 **speedup** 理想應該會和 **block** 大小成比，也就是 **16\*16** 會是 **8\*8** 四倍...以此類推，但實驗做出來明顯不是這樣，但整體 **speedup** 其實仍算有維持差不多的斜率。

2. 對 hw3-1 部分有花點時間驗證其 **strong scalability**，使用 **c20.1** 測 **edge:7594839 vertex:5000**，可以看到 **strong scalability** 因為有 **omp** 平行化 **Floyd Warshal** 外層兩個迴圈關係表現很好。



## Experience & Conclusion

本作業花很多時間在後面兩個 **cuda** 的部分上面，在第一部分，其實很多測資都是屬於 **edge** 稀疏的矩陣，理論上 **Floyd-Warshal** 時間複雜度  $O(n^3)$ ，並不適合 **edge** 稀疏的矩陣，可能用 **Johnson** 更適合，但

Johnson 的實作複雜且不易加入平行化，我最終還是先選擇 Floyd-Warshall+平行化通過 judge，想說之後再回來實驗 Johnson，結果後來就花了大量時間在後面兩個關於 cuda 的部分上，沒時間回來弄 Johnson 了。

不過在 cuda 上面我覺得學習 cuda 程式帶給我一種寫程式時的「空間感」，好像以前寫的程式都是線性的，以前平行化想的都只是如何把 for 迴圈拆分。如今突然變成二維甚至三維的感覺，進入另一個新世界，還學到 share memory、注意 bank conflict 概念等等，一口氣前進好多步。